

# Evolutionary Algorithms (EA)

Omar Mohammed

# Optimization

- Search for a solution to a defined problem
  - Unconstrained optimization
  - Constrained optimization
  - Satisfaction problem
  - Multiobjective problems
- Core of machine learning
  - Optimize the neural network weights (gradient descent)
  - Optimize the neural network hyperparameters (random / greedy)

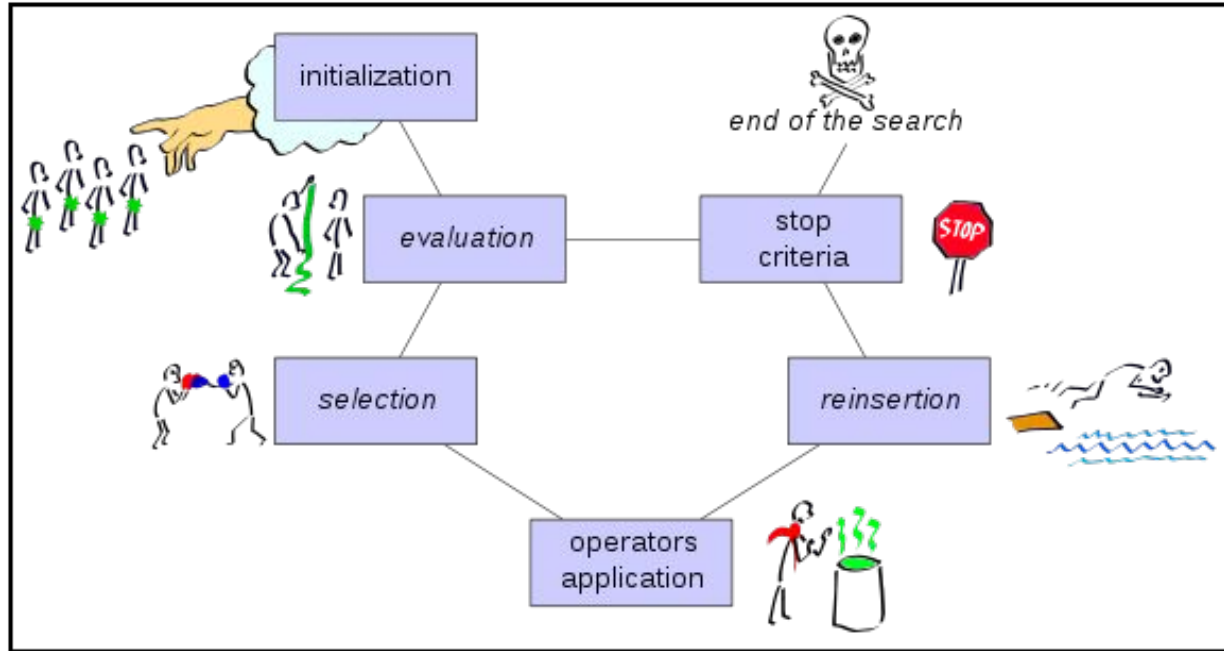
# Types of optimization

- Assumption based (more efficient, less general)
  - Convex optimization
  - Gradient descent
    - Convex/non-convex, but differentiable
- Agnostic (more general, less efficient)
  - Random search
  - Meta-heuristics
    - Evolutionary algorithms
    - Genetic programming
    - Ant-colony
    - Bee optimization
    - ....etc

# General Steps

1. (initialization) Start with N random solutions
  - a. Generation 0
  - b. Generation has N individuals
  - c. The group of individuals is called a population
2. (evaluation) Evaluate the quality of each solution
3. Are we there yet?
  - a. If not, continue
4. (secret sauce)
  - a. Select the potential solution
  - b. Little modification (mutation) and some sex (crossover)
  - c. New children are born
  - d. New generation is formed
5. Repeat

# What is an evolutionary algorithm?



*General schema of an Evolutionary Algorithm (EA)*


# Secret sauce

- Mutation

- Randomly change part of the individual
  - Ex: Add little random noise to the weights of a neural network
- Encourage exploitation

- Cross-over



- Exchange of genes
  - Between two individuals
- Encourages exploration
- Harder to design



- Most researchers stick to Mutation only

- Lots of work support this idea

# Simple genetic algorithm performance



# Example: simulated annealing

1. Initialize the system configuration.

Randomize  $\mathbf{x}(0)$ .

2. Initialize  $T$  with a large value.

3. **Repeat:**

- a. **Repeat:**

- i. Apply random perturbations to the state  $\mathbf{x} = \mathbf{x} + \Delta\mathbf{x}$ .

- ii. Evaluate  $\Delta E(\mathbf{x}) = E(\mathbf{x} + \Delta\mathbf{x}) - E(\mathbf{x})$ :

**if**  $\Delta E(\mathbf{x}) < 0$ , keep the new state;

**otherwise**, accept the new state with probability  $P = e^{-\frac{\Delta E}{T}}$ .

**until** the number of accepted transitions is below a threshold level.

- b. Set  $T = T - \Delta T$ .

**until**  $T$  is small enough.




# Example: simulated annealing



Example: shape/brain of the robot



# Example: evolve controller for robots



# Example: soccer team




# MNIST classifier

- 2 layer ConvNet
- ~11K parameters
- For EA algorithms
  - Population 101
  - Generations 300

Method	Train Set	Test Set
Adam (BackProp) Baseline	99.8	98.9
Simple GA	82.1	82.4
CMA-ES	98.4	98.1
OpenAI-ES	96.0	96.2
PEPG	98.5	98.0

# MNIST classifier



# Record in Atari games - simple genetic algorithm

Deep Neuroevolution: Genetic Algorithms are a Competitive Alternative for Training Deep Neural Networks for Reinforcement Learning

- Super simple
- State of the art performance in Many Atari games

---

## Algorithm 1 Simple Genetic Algorithm

---

**Input:** mutation power  $\sigma$ , population size  $N$ , number of selected individuals  $T$ , policy initialization routine  $\phi$ .

**for**  $g = 1, 2 \dots G$  generations **do**

**for**  $i = 1, \dots, N$  in next generation's population **do**

**if**  $g = 1$  **then**

$\mathcal{P}_i^g = \phi(\mathcal{N}(0, I))$  {initialize random DNN}

$F_i^g = F(\mathcal{P}_i^g)$  {assess its fitness}

**else**

**if**  $i = 1$  **then**

$\mathcal{P}_i^g = \mathcal{P}_i^{g-1}; F_i^g = F_i^{g-1}$  {copy the elite}

**else**

$k = \text{uniformRandom}(1, T)$  {select parent}

                Sample  $\epsilon \sim \mathcal{N}(0, I)$

$\mathcal{P}_i^g = \mathcal{P}_k^{g-1} + \sigma\epsilon$  {mutate parent}

$F_i^g = F(\mathcal{P}_i^g)$  {assess its fitness}

    Sort  $\mathcal{P}^g$  and  $F^g$  with descending order by  $F^g$

**Return:** highest performing policy,  $\mathcal{P}_1^g$

---

# Why to use EA?

- Fast prototyping
  - In NN for example: No need to think about differentiability
- Hard to describe the evaluation function
  - Relationship between speed and shape of robot

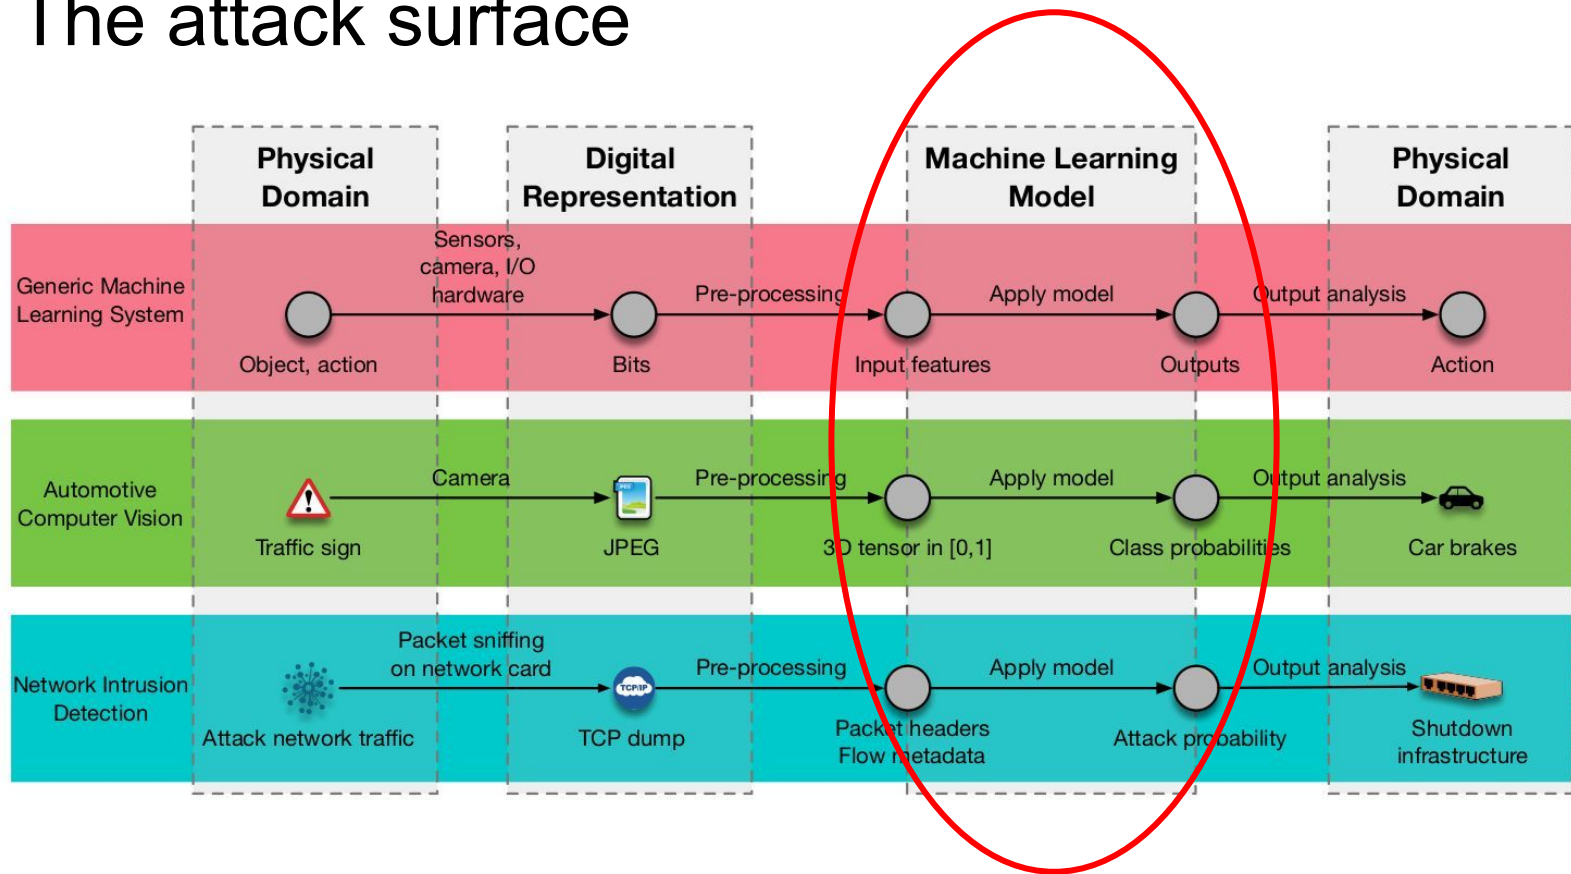


# Adversarial examples


# Adversarial examples

- Definition:
  - Data examples, that exploit the model limited knowledge about reality
- Important:
  - Compromise the integrity of the predictions (wrt expected outcome)
    - Can we trust the model results?
  - Compromise the availability of the system deploying machine learning
    - Effect of real-life data on the model


# The attack surface



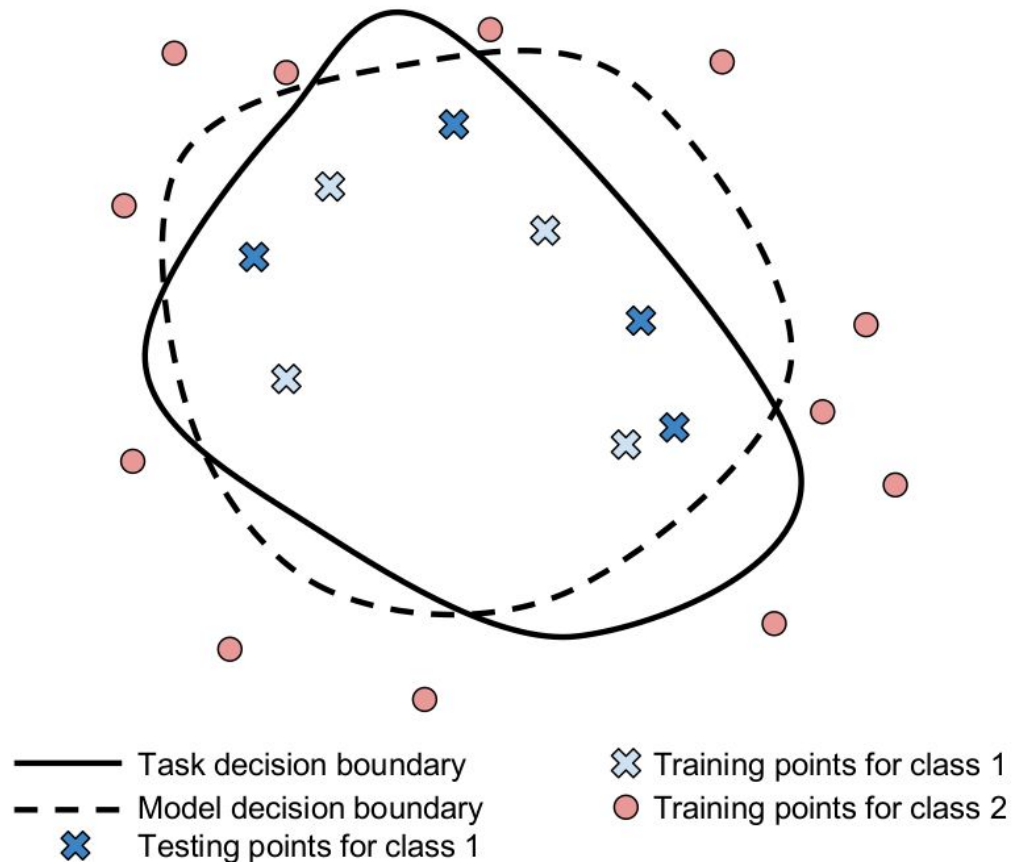
# Generalization error




# Generalization error



# Generalization error



# Generalization error




# Types of adversarial examples

- Training time
  - Ex: Training in the presence of malicious errors

- Inference time

- White-box attack (model inspection)



- Black-box attack (model query)





# Experiment setup

1. Train a model on an MNIST data set (the evaluator)

60K training data

10K test data

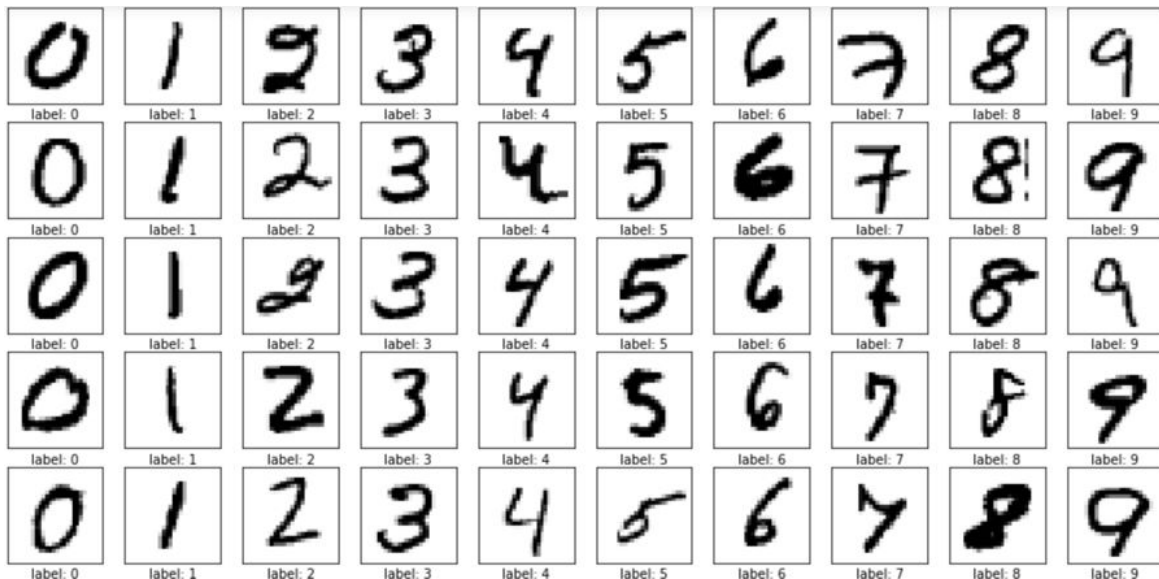
64 x 64 images

- a. For kNN, I use a subsample

1617 training data

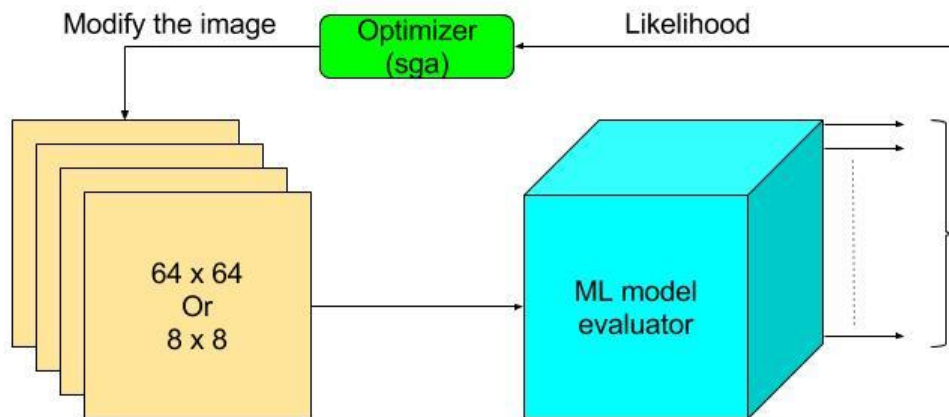
180 test data

8 x 8 images



# Experiment setup


## 2. Query the model (black box optimization)




# Logistic regression

- Train acc:  
92.7%
- Test acc:  
92%


Target: 0, Prob: 1.0




Target: 1, Prob: 1.0




Target: 2, Prob: 1.0




Target: 3, Prob: 1.0




Target: 4, Prob: 0.93




Target: 5, Prob: 1.0




Target: 6, Prob: 1.0




Target: 7, Prob: 1.0



Target: 8, Prob: 1.0




Target: 9, Prob: 1.0




# Multilayer perceptron

- 1 hidden layer, 50 neurons
- Train accuracy 98.57 %
- Test accuracy 97.1 %


Target: 0, Prob: 1.0




Target: 1, Prob: 1.0




Target: 2, Prob: 1.0




Target: 3, Prob: 1.0




Target: 4, Prob: 1.0




Target: 5, Prob: 1.0




Target: 6, Prob: 1.0




Target: 7, Prob: 1.0



Target: 8, Prob: 1.0




Target: 9, Prob: 1.0




# kNN - 5 neighbors

- Train acc:  
99 %
- Test acc:  
96.1 %


Target: 0, Prob: 1.0




Target: 1, Prob: 1.0




Target: 2, Prob: 1.0




Target: 3, Prob: 1.0




Target: 4, Prob: 1.0




Target: 5, Prob: 1.0




Target: 6, Prob: 1.0




Target: 7, Prob: 1.0



Target: 8, Prob: 1.0



Target: 9, Prob: 1.0






# Random Forest - 30 estimator - Same time


- Train acc:  
99 %
- Test acc:  
96.57%

**Harder to  
break...  
But  
breakable**


Target: 0, Prob: 0.73




Target: 1, Prob: 0.07




Target: 2, Prob: 0.87




Target: 3, Prob: 0.87




Target: 4, Prob: 0.47




Target: 5, Prob: 0.57




Target: 6, Prob: 0.63




Target: 7, Prob: 0.2



Target: 8, Prob: 0.97




Target: 9, Prob: 0.43



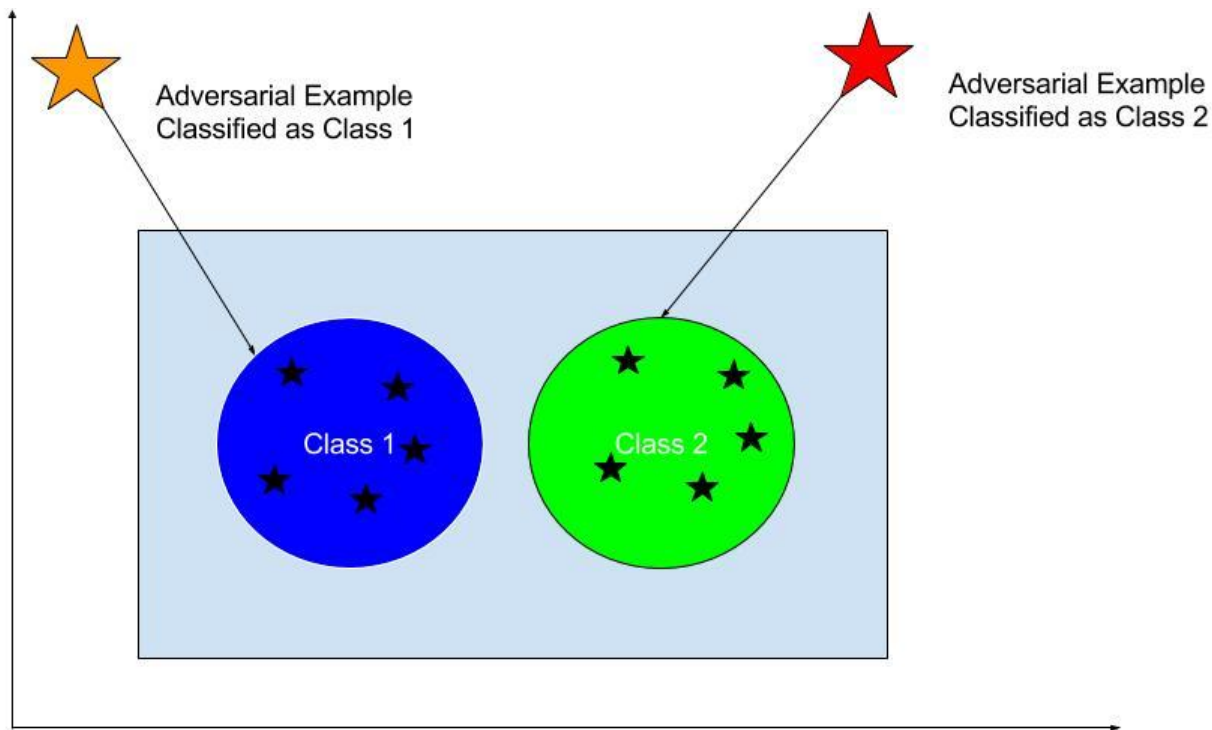
# Analysis on kNN performance

- Adversarial examples can be Very far from test examples
- Possible extra regularization



# Analysis on kNN performance


- What if we add an extra objective? (Regularization)
  - Max likelihood
  - Min the distance to nearest neighbors






# kNN: likelihood + distance to nearest neighbor

- Letters starting to appear
- Mount to an averaging Problem
- Target 1 is interesting!
  - Started from the position Closer to 8
  - A bad/unclear objective Can ruin the example



# Visualization for the kNN with the new loss function



# Resources

- <http://blog.otoro.net/2017/10/29/visual-evolution-strategies/>
- PyGMO optimization framework  
<https://esa.github.io/pagmo2/index.html>
  - Nice example  
[https://esa.github.io/pagmo2/docs/python/tutorials/coding\\_udp\\_simple.html](https://esa.github.io/pagmo2/docs/python/tutorials/coding_udp_simple.html)
- [https://www.tutorialspoint.com/genetic\\_algorithms/index.htm](https://www.tutorialspoint.com/genetic_algorithms/index.htm)